

1 Test-Driven Development (TDD)

Test-Driven Development [1, 4], abbreviated as TDD, is a technique that concerns the development *process*, rather than the structure of the developed *product*. Usually you cannot see in the end product whether TDD was applied or not. As the name Test-Driven Development suggests, testing plays a leading role. In particular, automated test cases are an integral part of delivered software. Furthermore, test cases are developed *before* product code is written. In fact, the development of the product is *driven* by the test cases. Most of the documentation is incorporated in the source code.

We explain **how** to apply TDD in the context of this course and **why** to apply it. We also present some **examples** and **exercises** using the Java programming language and the JUnit testing framework.

2 How: TDD steps

In TDD as applied in this course, development steps are performed in the following order. See [5] for the use of Mercurial to track this workflow.

1. Gather and analyze the requirements to realize, e.g. use cases or user stories.¹
2. Decide which requirement to develop next, and what module² it concerns.
3. Specify the module informally in the source code, using natural language.
4. Specify the module formally, by providing its *interface* and *contract* [2, 3].
5. Design and implement a test case³; document its motivation in the source code. This test case will now fail, lacking a module implementation.
6. Design and implement module features, intended to make the test case pass.
7. Test the new module implementation, and fix defects found (back to step 6).
8. Where relevant, improve the structure of the code (*refactoring*), preserving correctness by running all test cases again (*regression testing*).
9. Repeat from step 5, until the entire module has been implemented and tested.
10. If, after delivery, a defect is reported for the module *in vivo*, then, first, add a test case to detect that defect *in vitro*, and only then repair the defect.

It is recommended to write only **small fragments** of code at a time, and run all test cases after each change. This applies to both product code and test code.

¹In this course, requirements will be given to you.

²Typically, this involves a new class, but possibly a single method, or a package of related classes.

³The first time, you have to create a corresponding test module, to hold the automated test cases.

3 Why: Motivation for TDD

First, let's look at why one should *test* software *before* delivery.

- When humans produce software (or anything else) of some complexity, it is inevitable that they make mistakes and introduce **defects**. The presence of defects reduces the product's **quality**. The longer a defect remains undetected, the higher the **cost**⁴ of repairing it. Therefore, one should attempt to detect defects as soon as possible. Defects can be detected by *reviewing* intermediate artifacts, such as requirements documents and design documents, and also source code. But reviewing will not catch all defects. *Testing* is a complementary way of detecting defects based on code execution, applied after reviewing. Both reviewing and testing are necessary.

Why should one *write down* test cases, instead of improvising them on the fly?

- Test cases that are written down are better **verifiable** and **reproducible** than improvised test cases, provided that each test case clearly states
 - its motivation, i.e., its relationship to requirements, and reason for choosing this particular input combination;
 - which concrete input values to use;
 - which output values to consider;
 - precise criteria to decide between passing and failing.
- Test cases that are written down also serve as **documentation**, by making the requirements more concrete.
- The activity of creating test cases also helps to **verify the requirements**, viz. that those requirements are understandable.

Test cases can be performed manually, or their execution can be automated. Let's look at reasons for insisting on *automated* test cases as part of delivered software.

- **Manual testing is tedious, time consuming, and error prone.** Hence, it does not invite repeated application of the same test cases.
- **Automated test cases are more convenient, faster, and more reliable.**
 - They consist of **separate test code** that *invokes functionality* of the "subject under test", *captures the result*, and *reports about pass/fail*.
 - They offer the ability to (re)run all test cases, and report on the results.

⁴Experimental evidence shows that the cost of repair increases *exponentially* with time elapsed between injection of a defect and its subsequent detection.

N.B. Automated testing is (usually) not about automatically generating appropriate input values and code to verify output for correctness. Even with automated test cases, the developer will have to design the test cases, deciding on appropriate input values and how to check output for correctness. Some IDEs, such as NetBeans, do have the ability to automatically generate template test code, that must then be completed manually.

- It is easier to **inspect the results** of automated test cases.
- As documentation, automated test cases are an **operationalization** of the requirements by specifying, in an executable way,
 - how to invoke the tested feature;
 - how to use or observe the effect of the tested feature;
 - how to decide between passing and failing.

The motivation of a test case must be documented separately in a comment.

- Automated test cases serve as a client to the module under construction. Their development may reveal problems with the module's interface and contract early in the development cycle, before releasing it to clients in the production code. Thus, development of automated test cases helps **stabilize interfaces and contracts before releasing them**.
- When defects are present in code that was not exercised in a unit test, it can be hard to localize and repair them, leading to frustrating and unpredictable debugging⁵ work. Ad hoc debugging work is a waste of time, because it does not help you to do better in the future. And systematic debugging boils down to writing automated test cases anyway. So, why not write them in the first place? Good automated test cases **obviate later debugging**.
- Automated test cases can be rerun anytime; in particular, they can be rerun after changing the software (*regression testing*). They serve as a **safety net** when *refactoring* the code. That way, one can see to it that production code always works as expected.

Finally, let's look at why it is a good idea to write test code *before* product code. We already decided that test cases are needed, and that they must be automated in the form of program code. Choosing the order (before or after implementation) may seem like an irrelevant detail, but there are some important benefits to the TDD order.

- Because changes in a module's interface and contract lead to changes in its implementation, it is best to get the module's interface and contract in

⁵Testing concerns the *detection* of defects; debugging concerns the *repair* of detected defects.

good shape *before* starting to implement the module. As remarked earlier, automated test cases serve as a client of the module. Therefore, working on the automated test cases will confront you with the module's interface and contract. Development of automated test cases helps **stabilize interfaces and contracts *before* implementing them**. That is why it is a good idea to develop automated test cases *before* implementing the module.

- While developing automated test cases, you familiarize yourself with the module's interface and contract, and that will put you in a better position to fulfill the contract in an implementation. Development of automated test cases **encourages analysis of interfaces and contracts *before* implementing them**. This prevents you from rushing into implementation work unprepared.
- After completing (another part of) the module's implementation, the whole module must be tested. If the automated test cases are completed *before* the module's implementation, then there is no interruption between completing some implementation work and testing it. If a test case detects a defect, then you still have the implementation details fresh in your mind and can repair the defect quickly, and immediately retest it. Having automated test cases ready *before* implementing a module, allows you to **focus on implementing the module and getting it to work**, without being interrupted by test case development. This is also more satisfactory for a developer: once you have completed some implementation code, you can immediately try it out.

Exercises Answer the following questions in your own words.

1. In TDD, you implement the test cases for a module (cf. Step 5) before actually implementing that module (cf. Step 6). Explain why the workflow in Section 3 consists of more than just the two Steps 5 and 6.
2. What is another way of finding defects in source code besides testing it?
3. Why should software be tested before delivery?
4. Why is improvised testing not such a good idea?
5. Explain three benefits of writing down test cases.
6. Which four points do you need to address when defining a test case?
7. What is manual testing and what are its disadvantages?
8. How do automated test cases work?
9. What are additional benefits of automated test cases?

10. Explain how automated test cases help to shape the interface and contracts of a module.
11. Explain how automated test cases can simplify debugging.
12. Explain why it is useful that automated test cases can be rerun easily.
13. Explain three reasons why it is better to develop test cases *before* developing the implementation of a module.

4 TDD for a single function

To illustrate TDD in the small, we consider the development of

[Requirement] a function that counts the number of decimal digits in a non-negative integer.

We put the source code in the class `TDDForCountDigitsMethod`.

Give an **informal specification**, that is, provide a summary sentence in a *javadoc* comment, and a header (method name, parameter names and types, return type):

```
/**
 * Counts the decimal digits of a number.
 */
public static int countDigits(long n) {
}
```

Elaborate the informal specification, to remove ambiguities and imprecisions:

```
/**
 * Counts the decimal digits of a number.
 * This concerns a non-negative integer, assumed to be
 * written in decimal notation without leading zeroes.
 */
public static int countDigits(long n) {
}
```

Formalize the specification, that is, provide a **contract**⁶:

```
/**
 * Counts the decimal digits of a number.
 * This concerns a non-negative integer, assumed to be
 * written in decimal notation without leading zeroes.
 *
 * @param n the number whose digits are counted
```

⁶Here, we do not yet aim for *robustness*.

```

    * @return the number of decimal digits in {@code n}
    * @pre {@code 0 <= n}
    * @post {@code \result = (\min int k; 1 <= k; n < 10 ^ k)}
    */
    public static int countDigits(long n) {
    }

```

We commit our work, and next create class `TDDForCountDigitsMethodTest` for our test cases. We will use the JUnit testing framework, version 4.x. NetBeans puts the test class under **Test Packages**, and it can generate template test code, including these imports:

```

import static org.junit.Assert.*;
import org.junit.Test;

```

Because we are testing a **static** method, we do not need a class instance. Still, it is useful to have a short name for the class under test:

```

/** Subject Under Test. Only static members are used. */
private static final TDDForCountDigitsMethod SUT = null;

```

A first test case, with the smallest allowed number as input:

```

1  /**
2   * Test of countDigits method, of class TDDForCountDigitsMethod.
3   * Boundary case: smallest n
4   */
5  @Test
6  public void testCountDigits0() {
7      System.out.println("countDigits(0)");
8      long n = 0L;
9      int expectedResult = 1;
10     int result = SUT.countDigits(n);
11     assertEquals("result", expectedResult, result);
12 }

```

In JUnit, each test case is implemented in a parameterless **public void** method (line 6) annotated as `@Test` (line 5). The motivation is put in a javadoc comment (lines 1–4). The print statement (line 7) is helpful when inspecting test output. The test case first sets up the input (line 8), and prepares to check the output against the known expected result (line 9). Then, it calls the method to be tested with the input, and captures the result (line 10). Finally, it uses a method of the form `assertXxx(...)` to decide on pass/fail and report that decision to the JUnit framework (line 11). Each test case can include multiple `assertXxx` (and `fail`) calls, but it should focus on handling a single, independent, test case, because execution of a test case will be interrupted the moment an `assertXxx` fails (or `fail` is called). We commit our work.

Our first test case fails; in fact, the implementation does not even compile. But that is easy to fix:

```
public static int countDigits(long n) {
    return 0; // to make it compile
}
```

Not surprisingly, it still fails. Here is an implementation that is easy to understand, using a *Linear Search*:

```
public static int countDigits(long n) {
    // Linear Search for smallest k >= 1 with n < 10 ^ k
    int k = 1;
    long p = 10L;
    // invariant: 1 <= k && p == 10 ^ k
    while (! (n < p)) {
        p *= 10;
        ++ k;
    }
    // n < p == 10 ^ k
    return k;
}
```

It passes the first test case. Let's commit our work, and add another test case, with the largest number for the smallest result, i.e. $n = 9$. We refactor the test cases to reduce code duplication:

```
/**
 * Invokes countDigits(n) and checks result against expectation.
 *
 * @param n the number whose digits are counted
 * @param expResult the expected result
 * @pre {@code countDigits(n).pre}
 */
private void checkCountDigits(long n, int expResult) {
    System.out.println("countDigits(" + n + ")");
    int result = SUT.countDigits(n);
    assertEquals("result", expResult, result);
}

/**
 * Test of countDigits method, of class TDDForCountDigitsMethod.
 * Boundary case: largest n with smallest result
 */
@Test
public void testCountDigits9() {
```

```
        checkCountDigits(9L, 1);
    }
}
```

It passes. Add one more test case, the smallest number for result 2, i.e. $n = 10$.

```
/**
 * Test of countDigits method, of class TDDForCountDigitsMethod.
 * Boundary case: smallest n with result 2
 */
@Test
public void testCountDigits10() {
    checkCountDigits(10L, 2);
}
```

This also passes. Commit it. Time for a more systematic series of test cases, which includes all earlier test cases. All smallest and largest numbers with results up to 5:

```
/**
 * Test of countDigits method, of class TDDForCountDigitsMethod.
 * Boundary cases: smallest and largest numbers with small results
 */
@Test
public void testCountDigitsSmall() {
    long n = 1L;
    for (int r = 0; r <= 5; ++ r) {
        // n == 10 ^ r
        checkCountDigits(n - 1, Math.max(1, r));
        checkCountDigits(n, r + 1);
        n *= 10;
    }
}
```

All these tests pass. However, the implementation with a Linear Search has the risk of producing an overflow, because p is made to exceed n . So, let's add a test case for the largest possible value (the timeout is needed, because overflow⁷ may cause the Linear Search to loop endlessly):

```
/**
 * Test of countDigits method, of class TDDForCountDigitsMethod.
 * Boundary case: largest long value.
 * N.B. Overflow causes linear search to loop endlessly.
 */
@Test(timeout = 1000)
public void testCountDigitsMaxValue() {
    checkCountDigits(Long.MAX_VALUE, 19);
}
```

⁷Java will silently 'overflow' $p \ast = 10$ into the negative numbers.

This test case fails with a timeout. Commit it. Using the test cases as a safety net, we try to improve the implementation, by decreasing n , rather than increasing p :

```
public static int countDigits(long n) {
    int result = 1;
    while (n != 0) {
        n /= 10;
        ++ result;
    }
    return result;
}
```

This fails for $n > 0$, so the implementation is wrong (a *one-off*⁸ error). A fix:

```
public static int countDigits(long n) {
    int result = 1;
    while (10 <= n) {
        n /= 10;
        ++ result;
    }
    return result;
}
```

This passes all test cases, including the largest value, so we can have some confidence in the new implementation. Commit the latest work.

Exercises Apply TDD to address the following requirements.

14. Introduce a second parameter for the radix (base) r , which was hard coded as 10.
15. Make the function robust by checking that $0 \leq n$ and $2 \leq r$, throwing `IllegalArgumentException` if not. Also test for robustness.

5 TDD for a single class

Here is a summary of the steps for developing a class that will serve as an *Abstract Data Type*.⁸ Steps 1 and 2 are as before.

3. Specify the class **informally**:

- (a) provide a **one-sentence summary** of the class, in a *javadoc* comment;
- (b) choose an appropriate **class name**;

⁸We will no longer explicitly indicate when to commit.

- (c) when relevant, choose **generic type parameters** and their **constraints**;
 - (d) elaborate the informal description to obtain a more complete description of the class as a whole, from the perspective of a client of the class.
4. Specify the class **formally**, by providing
 - (a) a **model** defining the abstract set of values, in terms of specification variables, and/or basic queries and basic commands,
 - (b) **public invariants** constraining the model to intended values,
 - (c) **contracts** for its operations (constructors, queries, and commands), in terms of the model.

The result is a class without implementation: no data representation and empty method bodies. Consider putting this in an **abstract class** or **interface**. Note that not all operations have to be added at once.

5. Create a corresponding **unit test** class. Add a new **test case** for an operation.
6. Implement (part of) a **data representation** (or **rep** for short) in terms of **private** instance variables, and also provide a **rep(resentation) invariant** and **abstraction function**. Implement the operation of the preceding step.
7. Test the implemented operation.
8. Where relevant, refactor code to improve the structure.
9. Repeat from step 5 until done.

5.1 Example

To illustrate TDD of a class, we consider the development of

[Requirement] an Abstract Data Type to maintain a relation on small integers, in the range from 0 to a given upper bound.

Give an **informal specification**, that is, provide a summary sentence in a *javadoc* comment, and a header (class name; here, there are no generic type parameters):

```
/**
 * An {@code IntRelation} object maintains a relation on small integers.
 */
public abstract class IntRelation {
}
```

We make the class abstract, in order to provide multiple implementations. Elaborate the informal specification, to remove ambiguities and imprecisions:

```

/**
 * An {@code IntRelation} object maintains a relation on small integers.
 * The relation is a subset of {@code [0..n) x [0..n)},
 * where {@code n}, called the extent of the relation,
 * is given in the constructor and is immutable.
 * There are operations to test membership,
 * and to add and remove pairs.
 */
public abstract class IntRelation {
}

```

Formalize the specification⁹, that is, provide a **model** with **public invariants**, in the class javadoc:

```

 * Model: subset of {@code [0..extent()) x [0..extent())}
 *
 * @inv {@code NonNegativeExtent: 0 <= extent()}

```

And provide **contracts** for one or more operations. Here, we have decided to start with just the constructor, because that already gives us something to test and implement.

```

/**
 * Constructs an empty relation of given extent.
 *
 * @param n extent of the new relation
 * @pre {@code 0 <= n}
 * @post {@code this == [ ] && this.extent() == n}
 */
public IntRelation(final int n) {
}

```

To test concrete implementations of the abstract class, we define general test cases in an abstract test class that can be reused for each concrete implementation:

```

public abstract class IntRelationTestCases {

    /** Test fixture. */
    protected IntRelation instance;

    /**
     * Sets instance to a newly constructed relation of given extent.
     * (This is a kind of factory method; cf. Factory Method Design Pattern.)
     *
     * @param n the given extent
     * @pre {@code 0 <= n}
     * @modifies {@code instance}
     * @post {@code instance.extent() == n && AF(instance) = [ ]}
     */
}

```

⁹Again, we do not yet aim for robustness.

```

    protected abstract void setInstance(final int n);
}

```

A concrete test class must override `setInstance` (see below). Here is a test case for the constructor. Because there are not yet any queries to inspect the constructed object, we only test whether construction succeeds without throwing anything (`Error` or `Exception`).

```

/** Tests the constructor with small extent values. */
@Test
public void testConstructor() {
    System.out.println("constructor(int)");
    for (int n = 0; n <= 3; ++ n) {
        setInstance(n);
    }
}

```

5.1.1 Implementation with boolean arrays

The test cannot yet be applied, because we do not have a concrete implementation. We first work on an implementation that uses a two-dimensional array of `boolean` to store the *incidence matrix* of the relation. We provide a javadoc comment, class name, the representation, representation invariants, and abstraction function:

```

/**
 * An implementation of {@link IntRelation} using nested arrays.
 * For relations with a small extent this may work faster.
 * However, relations with a large extent use more memory,
 * even when they are sparse (have few related pairs).
 */
public class IntRelationArrays extends IntRelation {

    /** Representation of the relation. */
    protected final boolean[][] relation;

    /*
     * Representation invariants
     *
     * NotNull: relation != null
     * Extent: relation.length == extent()
     * ElementsNotNull: (\forall i; relation.has(i); relation[i] != null)
     * ElementsSameSize: (\forall i; relation.has(i);
     *     relation[i].length == relation.length)
     *
     * Abstraction function
     *
     * AF(this) = set of (a, b) such that relation[a][b] holds
     */
}

```

We did not forbid *sharing*, where `relation[i] == relation[j]` for `i != j`. It is not a problem for correctness of the constructor if all `relation[i]` equal the same all-**false** array (empty set). But do you see that it will make it harder—but not impossible—to implement `add`? We shall avoid sharing, although the rep invariant allows it.

Next we create a concrete test class for this implementation, overriding `setInstance`:

```
public class IntRelationArraysTest extends IntRelationTestCases {

    @Override
    protected void setInstance(final int n) {
        instance = new IntRelationArrays(n);
    }
}
```

When we run the test case, it fails, because we have not provided a constructor. So, we implement a constructor:

```
public IntRelationArrays(final int n) {
    super(n);
    relation = new boolean[n][n];
}
```

The test case passes.

Next, we intend to implement `isRepOk` to check the representation invariants. Note that developing this method early is valuable to prevent future debugging frustrations. If we would not do it now, and later would run into a problem that concerns the representation, then we would probably spend extra time doing some ad hoc inspections (using `println`), rather than improve the testing infrastructure.

We first provide an informal specification:

```
/**
 * Checks whether the representation invariants hold.
 */
public abstract boolean isRepOk() throws IllegalStateException;
```

and elaborate it into a formal specification:

```
/**
 * Checks whether the representation invariants hold.
 *
 * @return whether the representation invariants hold
 * @throws IllegalStateException if precondition violated
 * @pre representation invariants hold
 * @modifies None
 * @post {@code \result}
 */
public abstract boolean isRepOk() throws IllegalStateException;
```

There are two reasons to throw an exception when the rep invariants do not hold:

1. Just returning **false** is not very informative, whereas the exception can carry more detailed information about the violation.
2. Proceeding with normal operation makes no sense, when a rep invariant fails.

Now, we redefine the constructor test case to invoke `isRepOk`:

```
/** Tests the constructor with small extent values. */
@Test
public void testConstructor() {
    System.out.println("constructor(int)");
    for (int n = 0; n <= 3; ++ n) {
        setInstance(n);
        assertTrue("isRepOk()", instance.isRepOk());
    }
}
```

Of course, the test fails to compile, and we proceed to implement `isRepOk`

```
@Override
public boolean isRepOk() {
    // NotNull
    if (relation == null) {
        throw new IllegalStateException("relation == null");
    }
    for (int i = 0; i != extent(); ++ i) {
        // ElementsNotNull
        if (relation[i] == null) {
            throw new IllegalStateException(
                "relation[" + i + "] == null");
        }
        // ElementsSameSize
        if (relation[i].length != extent()) {
            throw new IllegalStateException(
                "relation[" + i + "].length != extent()");
        }
    }
    return true;
}
```

Fortunately, it passes. Let's move on to introduce a query to retrieve the extent. First, the informal specification:

```
/**
 * Returns the extent of this relation.
 */
public abstract int extent();
```

And then elaborated into a formal specification:

```
/**
 * Returns the extent of this relation.
```

```

    *
    * @return extent of this relation
    * @pre {@code true}
    * @modifies None
    * @post (basic query)
    */
    public abstract int extent();

```

Because this concerns a basic query, we cannot provide a formal postcondition.

Now, add a test case for it:

```

/** Tests the extent method with small relations. */
@Test
public void testExtent() {
    System.out.println("extent");
    for (int n = 0; n <= 3; ++ n) {
        setInstance(n);
        assertEquals("size", n, instance.extent());
        assertTrue("isRepOk()", instance.isRepOk());
    }
}

```

Again, it fails (to compile) for lack of an implementation. We fix that:

```

@Override
public int extent() {
    return relation.length;
}

```

Test again, ... with success (it is hard to go wrong this way).

In a similar way, we can, one by one, specify, provide a test for, and implement other queries and commands:

```

public boolean isValidPair(int a, int b) {
}

public abstract boolean areRelated(int a, int b);

public abstract void add(int a, int b);

public abstract void remove(int a, int b);

```

See the source code for details.

5.1.2 Implementation with a list of sets

A disadvantage of using arrays to store the relation is that it uses a lot of memory when the extent is larger. For *sparse* relations, which contain only relatively few pairs, another representation would be better. We briefly investigate an alternative implementation using a `List of Sets` from the Java Collection Framework.

First, we introduce a new class for the alternative implementation, with a javadoc comment, class name, representation, representation invariants, and abstraction function (the **import** section is not shown):

```
/**
 * An implementation of {@link IntRelation} using a {@link List} of {@link Sets}
 * For sparse relations with a large extent, this reduces memory usage.
 * However, there is a bit of performance overhead.
 */
public class IntRelationListOfSets extends IntRelation {

    /** Representation of the relation. */
    protected final List<Set<Integer>> relation;

    /**
     * Representation invariants
     *
     * NotNull: relation != null
     * SetsNotNull: (\forallall i; relation.has(i); relation.get(i) != null),
     *               where List.has(i) == 0 <= i < List.size()
     * SetsInExtent: (\forallall i; relation.has(i);
     *                relation.get(i) subset of [0 .. relation.size()))
     *
     * Abstraction function
     *
     * AF(this) = set of (a, b) such that relation.get(a).contains(b)
     */
}
```

We already have general test cases, and can reuse them again by overriding `setInstance`:

```
public class IntRelationListOfSetsTest extends IntRelationTestCases {

    @Override
    protected void setInstance(final int n) {
        instance = new IntRelationListOfSets(n);
    }
}
```

These tests all fail, for lack of implementations. All formal specifications are also already available. So, we can start implementing methods one-by-one, and rerun test cases after each method is done. For instance, here is the constructor:

```
public IntRelationListOfSets(final int n) {
    super(n);
    relation = new ArrayList<Set<Integer>>();
    for (int i = 0; i != n; ++ i) {
        relation.add(new HashSet<Integer>());
    }
}
```


Likewise, the remaining methods are implemented. See the source code for details. The Abstract Data Type is now ready for use by others. You see that it is easy to change the representation, once you have the safety net of a good set of test cases.

Exercises Apply TDD to address the following requirements.

16. Make `IntRelation` robust by checking preconditions of operations, and throwing `IllegalArgumentException` if they fail.

For each operation, first specify robustness in the abstract class, then develop a test case, and only then implement robustness in both concrete classes.

17. Provide a concrete class `IntRelationMapOfSets` extending `IntRelation`, based on a `Map of Sets`, to reduce memory usage further if an `a` is not related to any `b`. This can be formalized by the following representation invariant:

$$\text{relation.contains}(a) \iff (\exists b; ; (a, b) \text{ in } \text{AF}(\mathbf{this}))$$

where `AF` is the abstraction function. Methods `areRelated`, `add`, and `remove` need to take this into account. It must be robust.

18. Allow clients to use a for-each statement to iterate over all related pairs. That is, make `IntRelation` iterable, by (i) defining a (static nested) class `IntPair` for a record consisting of two integers `a` and `b`, and (ii) making `IntRelation` implement the interface `Iterable<IntPair>` to yield all related pairs one by one.

First specify the operation in the abstract class, then develop a test case, and only then implement it in both concrete classes.

19. Introduce an operation `relatedToFirst(int a)` returning an (object that implements the interface) `Iterable<Integer>` that iterates over all `b` such that `areRelated(a, b)` holds.

Introduce auxiliary method `isValid(a)` to express the precondition of the new operation. Also use it for `isValidPair` to reduce code duplication.

20. Generalize `IntRelation` to the generic interface `Relation<A, B>`, so that `Relation<Integer, Integer>` can be used as a relation on integers with unbounded extent. Provide one implementation `RelationMapOfSets<A, B>` using a `Map of Sets`. For iterating, introduce a generic record type `Pair<A, B>`.

References

- [1] Martin Fowler. *Test Driven Development*. Bliki, <http://martinfowler.com/bliki/TestDrivenDevelopment.html>

- [2] T. Verhoeff. *Notation*. A separate note for Programming Methods (2IPC0).
- [3] T. Verhoeff. *Specification*. A separate note for Programming Methods (2IPC0).
- [4] *Test-Driven Development*, Wikipedia, http://en.wikipedia.org/wiki/Test-driven_development
- [5] *Configuration Management*. A separate note for Programming Methods (2IPC0).

Author: Tom Verhoeff, Department of Mathematics & Computer Science, Eindhoven University of Technology